NeoSignal

AI Intelligence Daily · Multi-Source · Cross-Verified

10 March 2026

293 stories · 6 sources

293	7	286	0
STORIES	VERIFIED	CONFIRMED	EMERGING

How stories are selected

1 Scrape — HackerNews, Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Technology Review, The Verge AI, Wired AI, ArXiv CS.AI

2 Filter — 30+ AI keyword patterns (LLMs, safety, research, tooling, business)

3 Deduplicate — Title similarity matching across all sources (45%+ match = same story)

4 Score — Authenticity = base 0.5 + 0.3 per additional source + diversity bonus + HN score bonus

5 Rank — Highest authenticity first. Stories below 0.25 threshold are dropped.

Authenticity Tiers

- **VERIFIED (0.8–1.0)** Covered by 3+ independent sources including media outlets
- **CONFIRMED (0.5–0.79)** Covered by 2 sources or 1 high-quality media source
- **EMERGING (0.25–0.49)** Single-source but passes AI keyword filter

Source Breakdown

Source	Stories	Type
ArXiv CS.AI	253	Media
TechCrunch AI	16	Media
HackerNews	12	Community
VentureBeat AI	7	Media
MIT Tech Review	7	Media
The Verge AI	5	Media

VERIFIED INTELLIGENCE (7 stories)

1 **OpenAI and Google employees rush to Anthropic's defense in DOD lawsuit**

TechCrunch AI +1 other source VERI 80%

https://techcrunch.com/2026/03/09/openai-and-google-employees-rush-to-anthropics-defens...

2 **Anthropic launches code review tool to check flood of AI-generated code**

TechCrunch AI +1 other source VERI 80%

https://techcrunch.com/2026/03/09/anthropic-launches-code-review-tool-to-check-flood-of...

3 **OpenAI acquires Promptfoo to secure its AI agents**

TechCrunch AI +1 other source VERI 80%

https://techcrunch.com/2026/03/09/openai-acquires-promptfoo-to-secure-its-ai-agents/

4 **Anthropic sues Defense Department over supply-chain risk designation**

TechCrunch AI +1 other source VERI 80%

https://techcrunch.com/2026/03/09/anthropic-sues-defense-department-over-supply-chain-r...

5 **Will the Pentagon's Anthropic controversy scare startups away from defense work?**

TechCrunch AI +1 other source VERI 80%

https://techcrunch.com/2026/03/08/will-the-pentagons-anthropic-controversy-scare-startu...

6 **Anthropic's Pentagon deal is a cautionary tale for startups chasing federal contracts**

TechCrunch AI +1 other source VERI 80%

https://techcrunch.com/video/anthropics-pentagon-deal-is-a-cautionary-tale-for-startups...

7 **Online harassment is entering its AI era**

MIT Tech Review +1 other source VERI 80%

https://www.technologyreview.com/2026/03/05/1133962/online-harassment-is-entering-its-a...

CONFIRMED SIGNALS (286 stories)

8 **Show HN: Claude-replay – A video-like player for Claude Code sessions**

HN: 102

https://github.com/es617/claude-replay

9

Show HN: VS Code Agent Kanban: Task Management for the AI-Assisted Developer

HN: 90

https://www.appsoftware.com/blog/introducing-vs-code-agent-kanban-task-management-for-t...

10

Show HN: Kula – Lightweight, self-contained Linux server monitoring tool

HN: 90

https://github.com/c0m4r/kula

11

Show HN: 1v1 coding game that LLMs struggle with

HN: 26

https://yare.io

12

Show HN: Run 500B+ Parameter LLMs Locally on a Mac Mini

HN: 10

https://github.com/opengraviton/graviton

13

Show HN: SiClaw – Open-source AIOps with a hypothesis-driven diagnostic engine

HN: 9

https://github.com/scitix/siclaw

14

Show HN: Locode, a local first CLI that routes tasks to local LLMs or Claude

HN: 4

https://github.com/chocks/locode

15

Show HN: Finsight – A Privacy First, AI Credit Card and Bank Statement Analyzer

HN: 4

https://github.com/AJ/FinSight

16

Show HN: DopaLoop – Habit tracker for ADHD brains, local-first, no streaks

HN: 3

https://dopaloop.app/de

17

Show HN: I built a database of air fry times for 300+ foods

HN: 3

https://airfrypro.com

18

Show HN: Colchis Log – cryptographic audit trail for AI systems (Python)

HN: 3

https://github.com/GhurtSky-GR13/colchis-log

19

Show HN: ChatShell – 22MB AI Agent with 9 Built-In Tools (Tauri, Not Electron)

HN: 2

https://github.com/chatshellapp/chatshell-desktop

20

Owner of ICE detention facility sees big opportunity in AI man camps

https://techcrunch.com/2026/03/08/owner-of-ice-detention-facility-sees-big-opportunity-...

21

A roadmap for AI, if anyone will listen

https://techcrunch.com/2026/03/07/a-roadmap-for-ai-if-anyone-will-listen/

22

Google just gave Sundar Pichai a $692M pay package

https://techcrunch.com/2026/03/07/google-just-gave-sundar-pichai-a-692m-pay-package/

23

OpenAI hardware exec Caitlin Kalinowski quits in response to Pentagon deal

https://techcrunch.com/2026/03/07/openai-robotics-lead-caitlin-kalinowski-quits-in-resp...

24

OpenAI delays ChatGPT’s ‘adult mode’ again

https://techcrunch.com/2026/03/07/openai-delays-chatgpts-adult-mode-again/

25

Anthropic’s Claude found 22 vulnerabilities in Firefox over two weeks

https://techcrunch.com/2026/03/06/anthropics-claude-found-22-vulnerabilities-in-firefox...

26

Anthropic vs. the Pentagon, the SaaSpocalypse, and why competitions is good, actually

https://techcrunch.com/podcast/anthropic-vs-the-pentagon-the-saaspocalypse-and-why-comp...

27

Claude’s consumer growth surge continues after Pentagon deal debacle

https://techcrunch.com/2026/03/06/claudes-consumer-growth-surge-continues-after-pentago...

28

City Detect, which uses AI to help cities stay safe and clean, raises $13M Series A

https://techcrunch.com/2026/03/06/city-detect-uses-ai-to-help-cities-stay-safe-and-clean/

29

After Europe, WhatsApp will let rival AI companies offer chatbots in Brazil

https://techcrunch.com/2026/03/06/after-europe-whatsapp-will-let-rival-ai-companies-off...

30

Railway secures $100 million to challenge AWS with AI-native cloud infrastructure

https://venturebeat.com/infrastructure/railway-secures-usd100-million-to-challenge-aws-...

31

Claude Code costs up to $200 a month. Goose does the same thing for free.

https://venturebeat.com/infrastructure/claude-code-costs-up-to-usd200-a-month-goose-doe...

32

Listen Labs raises $69M after viral billboard hiring stunt to scale AI customer interviews

https://venturebeat.com/technology/listen-labs-raises-usd69m-after-viral-billboard-hiri...

33

Salesforce rolls out new Slackbot AI agent as it battles Microsoft and Google in workplace AI

https://venturebeat.com/technology/salesforce-rolls-out-new-slackbot-ai-agent-as-it-bat...

34

Nous Research's NousCoder-14B is an open-source coding model landing right in the Claude Code moment

https://venturebeat.com/technology/nous-researchs-nouscoder-14b-is-an-open-source-codin...

35

The creator of Claude Code just revealed his workflow, and developers are losing their minds

https://venturebeat.com/technology/the-creator-of-claude-code-just-revealed-his-workflo...

36

How AI is turning the Iran conflict into theater

https://www.technologyreview.com/2026/03/09/1134063/how-ai-is-turning-the-iran-conflict...

37

The Download: murky AI surveillance laws, and the White House cracks down on defiant labs

https://www.technologyreview.com/2026/03/09/1134050/the-download-ai-surveillance-laws-w...

38

The Download: 10 things that matter in AI, plus Anthropic’s plan to sue the Pentagon

https://www.technologyreview.com/2026/03/06/1133989/the-download-10-things-that-matter-...

39

The Download: an AI agent’s hit piece, and preventing lightning

https://www.technologyreview.com/2026/03/05/1133968/the-download-ai-agent-hit-piece-pre...

40

Bridging the operational AI gap

https://www.technologyreview.com/2026/03/04/1133642/bridging-the-operational-ai-gap/

41

The AI Doc is an overwrought hype piece for doomers and accelerationists alike

https://www.theverge.com/entertainment/890806/the-ai-doc-or-how-i-became-an-apocaloptim...

42

Netflix is buying Ben Affleck’s AI startup

https://www.theverge.com/streaming/889973/netflix-ben-affleck-interpositive-ai

43

Autonomous AI Agents for Option Hedging: Enhancing Financial Stability through Shortfall Aware Reinforcement Learning

https://arxiv.org/abs/2603.06587

44

Scaling Strategy, Not Compute: A Stand-Alone, Open-Source StarCraft II Benchmark for Accessible Reinforcement Learnin...

https://arxiv.org/abs/2603.06608

45

MultiGen: Level-Design for Editable Multiplayer Worlds in Diffusion Game Engines

https://arxiv.org/abs/2603.06679

46

Best-of-Tails: Bridging Optimism and Pessimism in Inference-Time Alignment

https://arxiv.org/abs/2603.06797

47

Making AI Evaluation Deployment Relevant Through Context Specification

https://arxiv.org/abs/2603.06811

48

Symmetry-Constrained Language-Guided Program Synthesis for Discovering Governing Equations from Noisy and Partial Obs...

https://arxiv.org/abs/2603.06869

49

LieCraft: A Multi-Agent Framework for Evaluating Deceptive Capabilities in Language Models

https://arxiv.org/abs/2603.06874

50

Not Too Short, Not Too Long: How LLM Response Length Shapes People's Critical Thinking in Error Detection

https://arxiv.org/abs/2603.06878

51

Enhancing the Detection of Coronary Artery Disease Using Machine Learning

https://arxiv.org/abs/2603.06888

52

Self-Supervised Multi-Modal World Model with 4D Space-Time Embedding

https://arxiv.org/abs/2603.07039

53

Animating Petascale Time-varying Data on Commodity Hardware with LLM-assisted Scripting

https://arxiv.org/abs/2603.07053

54

Grounding Machine Creativity in Game Design Knowledge Representations: Empirical Probing of LLM-Based Executable Synt...

https://arxiv.org/abs/2603.07101

55

Vision Language Models Cannot Reason About Physical Transformation

https://arxiv.org/abs/2603.07109

56

Improving reasoning at inference time via uncertainty minimisation

https://arxiv.org/abs/2603.07159

57

$\textbf{Re}^{2}$: Unlocking LLM Reasoning via Reinforcement Learning with Re-solving

https://arxiv.org/abs/2603.07197

58

A Cortically Inspired Architecture for Modular Perceptual AI

https://arxiv.org/abs/2603.07295

59

Shutdown Safety Valves for Advanced AI

https://arxiv.org/abs/2603.07315

60

FinSheet-Bench: From Simple Lookups to Complex Reasoning, Where LLMs Break on Financial Spreadsheets

https://arxiv.org/abs/2603.07316

61

The Third Ambition: Artificial Intelligence and the Science of Human Behavior

https://arxiv.org/abs/2603.07329

62

The Yerkes-Dodson Curve for AI Agents: Emergent Cooperation Under Environmental Pressure in Multi-Agent LLM Simulations

https://arxiv.org/abs/2603.07360

63

SoK: Agentic Retrieval-Augmented Generation (RAG): Taxonomy, Architectures, Evaluation, and Research Directions

https://arxiv.org/abs/2603.07379

64

AutoControl Arena: Synthesizing Executable Test Environments for Frontier AI Risk Evaluation

https://arxiv.org/abs/2603.07427

65

Do Machines Fail Like Humans? A Human-Centred Out-of-Distribution Spectrum for Mapping Error Alignment

https://arxiv.org/abs/2603.07462

66

COOL-MC: Verifying and Explaining RL Policies for Multi-bridge Network Maintenance

https://arxiv.org/abs/2603.07546

67

Memory for Autonomous LLM Agents:Mechanisms, Evaluation, and Emerging Frontiers

https://arxiv.org/abs/2603.07670

68

Rigidity in LLM Bandits with Implications for Human-AI Dyads

https://arxiv.org/abs/2603.07717

69

A Novel Multi-Agent Architecture to Reduce Hallucinations of Large Language Models in Multi-Step Structural Modeling

https://arxiv.org/abs/2603.07728

70

Large Language Model for Discrete Optimization Problems: Evaluation and Step-by-step Reasoning

https://arxiv.org/abs/2603.07733

71

SynPlanResearch-R1: Encouraging Tool Exploration for Deep Research with Synthetic Plans

https://arxiv.org/abs/2603.07853

72

SMGI: A Structural Theory of General Artificial Intelligence

https://arxiv.org/abs/2603.07896

73

EveryQuery: Zero-Shot Clinical Prediction via Task-Conditioned Pretraining over Electronic Health Records

https://arxiv.org/abs/2603.07900

74

Ares: Adaptive Reasoning Effort Selection for Efficient LLM Agents

https://arxiv.org/abs/2603.07915

75

Rel-MOSS: Towards Imbalanced Relational Deep Learning on Relational Databases

https://arxiv.org/abs/2603.07916

76

Adaptive Collaboration with Humans: Metacognitive Policy Optimization for Multi-Agent LLMs with Continual Learning

https://arxiv.org/abs/2603.07972

77

CMMR-VLN: Vision-and-Language Navigation via Continual Multimodal Memory Retrieval

https://arxiv.org/abs/2603.07997

78

S2S-FDD: Bridging Industrial Time Series and Natural Language for Explainable Zero-shot Fault Diagnosis

https://arxiv.org/abs/2603.08048

79

In-Context Reinforcement Learning for Tool Use in Large Language Models

https://arxiv.org/abs/2603.08068

80

Evidence-Driven Reasoning for Industrial Maintenance Using Heterogeneous Data

https://arxiv.org/abs/2603.08171

81

The Struggle Between Continuation and Refusal: A Mechanistic Analysis of the Continuation-Triggered Jailbreak in LLMs

https://arxiv.org/abs/2603.08234

82

Deconstructing Multimodal Mathematical Reasoning: Towards a Unified Perception-Alignment-Reasoning Paradigm

https://arxiv.org/abs/2603.08291

83

A Hierarchical Error-Corrective Graph Framework for Autonomous Agents with LLM-Based Action Generation

https://arxiv.org/abs/2603.08388

84

Efficient Policy Learning with Hybrid Evaluation-Based Genetic Programming for Uncertain Agile Earth Observation Sate...

https://arxiv.org/abs/2603.08447

85

OfficeQA Pro: An Enterprise Benchmark for End-to-End Grounded Reasoning

https://arxiv.org/abs/2603.08655

86

A Multi-Objective Optimization Approach for Sustainable AI-Driven Entrepreneurship in Resilient Economies

https://arxiv.org/abs/2603.08692

87

Evaluating Financial Intelligence in Large Language Models: Benchmarking SuperInvesting AI with LLM Engines

https://arxiv.org/abs/2603.08704

88

Agentic Critical Training

https://arxiv.org/abs/2603.08706

89

Deep Learning-Based Approach for Automatic 2D and 3D MRI Segmentation of Gliomas

https://arxiv.org/abs/2502.19760

90

A Robust Incomplete Multimodal Low-Rank Adaptation Approach for Emotion Recognition

https://arxiv.org/abs/2507.11202

91

Agentic SPARQL: Evaluating SPARQL-MCP-powered Intelligent Agents on the Federated KGQA Benchmark

https://arxiv.org/abs/2603.06582

92

A Coin Flip for Safety: LLM Judges Fail to Reliably Measure Adversarial Robustness

https://arxiv.org/abs/2603.06594

93

Distributionally Robust Geometric Joint Chance-Constrained Optimization: Neurodynamic Approaches

https://arxiv.org/abs/2603.06597

94

Building the ethical AI framework of the future: from philosophy to practice

https://arxiv.org/abs/2603.06599

95

Multi-Agent DRL for V2X Resource Allocation: Disentangling Challenges and Benchmarking Solutions

https://arxiv.org/abs/2603.06607

96

Consensus is Not Verification: Why Crowd Wisdom Strategies Fail for LLM Truthfulness

https://arxiv.org/abs/2603.06612

97

Annealed Co-Generation: Disentangling Variables via Progressive Pairwise Modeling

https://arxiv.org/abs/2603.06615

98

Not all tokens are needed(NAT): token efficient reinforcement learning

https://arxiv.org/abs/2603.06619

99

Pavement Missing Condition Data Imputation through Collective Learning-Based Graph Neural Networks

https://arxiv.org/abs/2603.06625

100

Grouter: Decoupling Routing from Representation for Accelerated MoE Training

https://arxiv.org/abs/2603.06626

101

Photons = Tokens: The Physics of AI and the Economics of Knowledge

https://arxiv.org/abs/2603.06630

102

SmartBench: Evaluating LLMs in Smart Homes with Anomalous Device States and Behavioral Contexts

https://arxiv.org/abs/2603.06636

103

SR-TTT: Surprisal-Aware Residual Test-Time Training

https://arxiv.org/abs/2603.06642

104

Performance Comparison of IBN orchestration using LLM and SLMs

https://arxiv.org/abs/2603.06647

105

HURRI-GAN: A Novel Approach for Hurricane Bias-Correction Beyond Gauge Stations using Generative Adversarial Networks

https://arxiv.org/abs/2603.06649

106

PaLMR: Towards Faithful Visual Reasoning via Multimodal Process Alignment

https://arxiv.org/abs/2603.06652

107

Science Literacy: Generative AI as Enabler of Coherence in the Teaching, Learning, and Assessment of Scientific Knowl...

https://arxiv.org/abs/2603.06659

108

Graph-of-Mark: Promote Spatial Reasoning in Multimodal Language Models with Graph-Based Visual Prompting

https://arxiv.org/abs/2603.06663

109

Accelerating Video Generation Inference with Sequential-Parallel 3D Positional Encoding Using a Global Time Index

https://arxiv.org/abs/2603.06664

110

Better Eyes, Better Thoughts: Why Vision Chain-of-Thought Fails in Medicine

https://arxiv.org/abs/2603.06665

111

Hybrid Orchestration of Edge AI and Microservices via Graph-based Self-Imitation Learning

https://arxiv.org/abs/2603.06669

112

calibfusion: Transformer-Based Differentiable Calibration for Radar-Camera Fusion Detection in Water-Surface Environm...

https://arxiv.org/abs/2603.06670

113

Does Semantic Noise Initialization Transfer from Images to Videos? A Paired Diagnostic Study

https://arxiv.org/abs/2603.06672

114

XAI and Few-shot-based Hybrid Classification Model for Plant Leaf Disease Prognosis

https://arxiv.org/abs/2603.06676

115

UWPD: A General Paradigm for Invisible Watermark Detection Agnostic to Embedding Algorithms

https://arxiv.org/abs/2603.06723

116

Regression Models Meet Foundation Models: A Hybrid-AI Approach to Practical Electricity Price Forecasting

https://arxiv.org/abs/2603.06726

117

Safe Transformer: An Explicit Safety Bit For Interpretable And Controllable Alignment

https://arxiv.org/abs/2603.06727

118

Calibrated Credit Intelligence: Shift-Robust and Fair Risk Scoring with Bayesian Uncertainty and Gradient Boosting

https://arxiv.org/abs/2603.06733

119

Rank-Factorized Implicit Neural Bias: Scaling Super-Resolution Transformer with FlashAttention

https://arxiv.org/abs/2603.06738

120

Heterogeneous Decentralized Diffusion Models

https://arxiv.org/abs/2603.06741

121

Enhancing Instruction Following of LLMs via Activation Steering with Dynamic Rejection

https://arxiv.org/abs/2603.06745

122

ButterflyViT: 354$\times$ Expert Compression for Edge Vision Transformers

https://arxiv.org/abs/2603.06746

123

Robotic Foundation Models for Industrial Control: A Comprehensive Survey and Readiness Assessment Framework

https://arxiv.org/abs/2603.06749

124

XMACNet: An Explainable Lightweight Attention based CNN with Multi Modal Fusion for Chili Disease Classification

https://arxiv.org/abs/2603.06750

125

Enhancing SHAP Explainability for Diagnostic and Prognostic ML Models in Alzheimer Disease

https://arxiv.org/abs/2603.06758

126

Diversity-Aware Adaptive Collocation for Physics-Informed Neural Networks via Sparse QUBO Optimization and Hybrid Cor...

https://arxiv.org/abs/2603.06761

127

SpatialMAGIC: A Hybrid Framework Integrating Graph Diffusion and Spatial Attention for Spatial Transcriptomics Imputa...

https://arxiv.org/abs/2603.06780

128

xaitimesynth: A Python Package for Evaluating Attribution Methods for Time Series with Synthetic Ground Truth

https://arxiv.org/abs/2603.06781

129

Physics-Informed Diffusion Model for Generating Synthetic Extreme Rare Weather Events Data

https://arxiv.org/abs/2603.06782

130

A Hybrid Machine Learning Model for Cerebral Palsy Detection

https://arxiv.org/abs/2603.06803

131

AI-Assisted Curation of Conference Scholarship: Compiling, Structuring, and Analyzing Two Decades of Presentations at...

https://arxiv.org/abs/2603.06814

132

"Dark Triad" Model Organisms of Misalignment: Narrow Fine-Tuning Mirrors Human Antisocial Behavior

https://arxiv.org/abs/2603.06816

133

Step-Level Visual Grounding Faithfulness Predicts Out-of-Distribution Generalization in Long-Horizon Vision-Language ...

https://arxiv.org/abs/2603.06828

134

Supporting Artifact Evaluation with LLMs: A Study with Published Security Research Papers

https://arxiv.org/abs/2603.06862

135

Physics-informed AI Accelerated Retention Analysis of Ferroelectric Vertical NAND: From Day-Scale TCAD to Second-Scal...

https://arxiv.org/abs/2603.06881

136

Post-Training with Policy Gradients: Optimality and the Base Model Barrier

https://arxiv.org/abs/2603.06957

137

Diffusion Controller: Framework, Algorithms and Parameterization

https://arxiv.org/abs/2603.06981

138

Masked Unfairness: Hiding Causality within Zero ATE

https://arxiv.org/abs/2603.06984

139

Foundational World Models Accurately Detect Bimanual Manipulator Failures

https://arxiv.org/abs/2603.06987

140

SuperSkillsStack: Agency, Domain Knowledge, Imagination, and Taste in Human-AI Design Education

https://arxiv.org/abs/2603.07016

141

RESCHED: Rethinking Flexible Job Shop Scheduling from a Transformer-based Architecture with Simplified States

https://arxiv.org/abs/2603.07020

142

MedSteer: Counterfactual Endoscopic Synthesis via Training-Free Activation Steering

https://arxiv.org/abs/2603.07066

143

Enhancing Consistency of Werewolf AI through Dialogue Summarization and Persona Information

https://arxiv.org/abs/2603.07111

144

Deep Expert Injection for Anchoring Retinal VLMs with Domain-Specific Knowledge

https://arxiv.org/abs/2603.07131

145

Emotion Transcription in Conversation: A Benchmark for Capturing Subtle and Complex Emotional States through Natural ...

https://arxiv.org/abs/2603.07138

146

Fine-Grained Table Retrieval Through the Lens of Complex Queries

https://arxiv.org/abs/2603.07146

147

From State Changes to Creative Decisions: Documenting and Interpreting Traces Across Creative Domains

https://arxiv.org/abs/2603.07184

148

A Miniature Brain Transformer: Thalamic Gating, Hippocampal Lateralization, Amygdaloid Salience, and Prefrontal Worki...

https://arxiv.org/abs/2603.07217

149

A Hybrid LTR-based System via Social Context Embedding for Recommending Solutions of Software Bugs in Developer Commu...

https://arxiv.org/abs/2603.07229

150

Do Deployment Constraints Make LLMs Hallucinate Citations? An Empirical Study across Four Models and Five Prompting R...

https://arxiv.org/abs/2603.07287

151

Adversarial Latent-State Training for Robust Policies in Partially Observable Domains

https://arxiv.org/abs/2603.07313

152

Agora: Teaching the Skill of Consensus-Finding with AI Personas Grounded in Human Voice

https://arxiv.org/abs/2603.07339

153

ConfHit: Conformal Generative Design with Oracle Free Guarantees

https://arxiv.org/abs/2603.07371

154

Backdoor4Good: Benchmarking Beneficial Uses of Backdoors in LLMs

https://arxiv.org/abs/2603.07452

155

"Better Ask for Forgiveness than Permission": Practices and Policies of AI Disclosure in Freelance Work

https://arxiv.org/abs/2603.07459

156

Where Do LLM-based Systems Break? A System-Level Security Framework for Risk Assessment and Treatment

https://arxiv.org/abs/2603.07460

157

Contact-Guided 3D Genome Structure Generation of E. coli via Diffusion Transformers

https://arxiv.org/abs/2603.07472

158

Give Them an Inch and They Will Take a Mile:Understanding and Measuring Caller Identity Confusion in MCP-Based AI Sys...

https://arxiv.org/abs/2603.07473

159

Interpretable-by-Design Transformers via Architectural Stream Independence

https://arxiv.org/abs/2603.07482

160

From Thinker to Society: Security in Hierarchical Autonomy Evolution of AI Agents

https://arxiv.org/abs/2603.07496

161

SketchGraphNet: A Memory-Efficient Hybrid Graph Transformer for Large-Scale Sketch Corpora Recognition

https://arxiv.org/abs/2603.07521

162

How Long Can Unified Multimodal Models Generate Images Reliably? Taming Long-Horizon Interleaved Image Generation via...

https://arxiv.org/abs/2603.07540

163

Nw\=ach\=a Mun\=a: A Devanagari Speech Corpus and Proximal Transfer Benchmark for Nepal Bhasha ASR

https://arxiv.org/abs/2603.07554

164

GRD-Net: Generative-Reconstructive-Discriminative Anomaly Detection with Region of Interest Attention Module

https://arxiv.org/abs/2603.07566

165

Evaluating Synthetic Data for Baggage Trolley Detection in Airport Logistics

https://arxiv.org/abs/2603.07645

166

AI-Driven Phase Identification from X-ray Hyperspectral Imaging of cycled Na-ion Cathode Materials

https://arxiv.org/abs/2603.07666

167

Compressed-Domain-Aware Online Video Super-Resolution

https://arxiv.org/abs/2603.07694

168

TDM-R1: Reinforcing Few-Step Diffusion Models with Non-Differentiable Reward

https://arxiv.org/abs/2603.07700

169

YAQIN: Culturally Sensitive, Agentic AI for Mental Healthcare Support Among Muslim Women in the UK

https://arxiv.org/abs/2603.07709

170

DECADE: A Temporally-Consistent Unsupervised Diffusion Model for Enhanced Rb-82 Dynamic Cardiac PET Image Denoising

https://arxiv.org/abs/2603.07759

171

QuadAI at SemEval-2026 Task 3: Ensemble Learning of Hybrid RoBERTa and LLMs for Dimensional Aspect-Based Sentiment An...

https://arxiv.org/abs/2603.07766

172

Learning embeddings of non-linear PDEs: the Burgers' equation

https://arxiv.org/abs/2603.07812

173

AI Misuse in Education Is a Measurement Problem: Toward a Learning Visibility Framework

https://arxiv.org/abs/2603.07834

174

DistillGuard: Evaluating Defenses Against LLM Knowledge Distillation

https://arxiv.org/abs/2603.07835

175

Slumbering to Precision: Enhancing Artificial Neural Network Calibration Through Sleep-like Processes

https://arxiv.org/abs/2603.07867

176

CCR-Bench: A Comprehensive Benchmark for Evaluating LLMs on Complex Constraints, Control Flows, and Real-World Cases

https://arxiv.org/abs/2603.07886

177

Reject, Resample, Repeat: Understanding Parallel Reasoning in Language Model Inference

https://arxiv.org/abs/2603.07887

178

Designing probabilistic AI monsoon forecasts to inform agricultural decision-making

https://arxiv.org/abs/2603.07893

179

SWE-Fuse: Empowering Software Agents via Issue-free Trajectory Learning and Entropy-aware RLVR Training

https://arxiv.org/abs/2603.07927

180

ELLMob: Event-Driven Human Mobility Generation with Self-Aligned LLM Framework

https://arxiv.org/abs/2603.07946

181

Aero-Promptness: Drag-Aware Aerodynamic Manipulability for Propeller-driven Vehicles

https://arxiv.org/abs/2603.07998

182

FedMomentum: Preserving LoRA Training Momentum in Federated Fine-Tuning

https://arxiv.org/abs/2603.08014

183

Alignment--Process--Outcome: Rethinking How AIs and Humans Collaborate

https://arxiv.org/abs/2603.08017

184

DyLLM: Efficient Diffusion LLM Inference via Saliency-based Token Selection and Partial Attention

https://arxiv.org/abs/2603.08026

185

Solution to the 10th ABAW Expression Recognition Challenge: A Robust Multimodal Framework with Safe Cross-Attention a...

https://arxiv.org/abs/2603.08034

186

DSH-Bench: A Difficulty- and Scenario-Aware Benchmark with Hierarchical Subject Taxonomy for Subject-Driven Text-to-I...

https://arxiv.org/abs/2603.08090

187

DC-W2S: Dual-Consensus Weak-to-Strong Training for Reliable Process Reward Modeling in Biological Reasoning

https://arxiv.org/abs/2603.08095

188

SaiVLA-0: Cerebrum--Pons--Cerebellum Tripartite Architecture for Compute-Aware Vision-Language-Action

https://arxiv.org/abs/2603.08124

189

Foley-Flow: Coordinated Video-to-Audio Generation with Masked Audio-Visual Alignment and Dynamic Conditional Flows

https://arxiv.org/abs/2603.08126

190

An explainable hybrid deep learning-enabled intelligent fault detection and diagnosis approach for automotive softwar...

https://arxiv.org/abs/2603.08165

191

TildeOpen LLM: Leveraging Curriculum Learning to Achieve Equitable Language Representation

https://arxiv.org/abs/2603.08182

192

MM-TS: Multi-Modal Temperature and Margin Schedules for Contrastive Learning with Long-Tail Data

https://arxiv.org/abs/2603.08202

193

Distributional Regression with Tabular Foundation Models: Evaluating Probabilistic Predictions via Proper Scoring Rules

https://arxiv.org/abs/2603.08206

194

Alignment-Aware and Reliability-Gated Multimodal Fusion for Unmanned Aerial Vehicle Detection Across Heterogeneous Th...

https://arxiv.org/abs/2603.08208

195

Exploring Deep Learning and Ultra-Widefield Imaging for Diabetic Retinopathy and Macular Edema

https://arxiv.org/abs/2603.08235

196

SCL-GNN: Towards Generalizable Graph Neural Networks via Spurious Correlation Learning

https://arxiv.org/abs/2603.08270

197

AdaCultureSafe: Adaptive Cultural Safety Grounded by Cultural Knowledge in Large Language Models

https://arxiv.org/abs/2603.08275

198

Evaluating LLM-Based Grant Proposal Review via Structured Perturbations

https://arxiv.org/abs/2603.08281

199

A Blockchain-based Traceability System for AI-Driven Engine Blade Inspection

https://arxiv.org/abs/2603.08288

200

Graph-Instructed Neural Networks for parametric problems with varying boundary conditions

https://arxiv.org/abs/2603.08304

201

Concept-Guided Fine-Tuning: Steering ViTs away from Spurious Correlations to Improve Robustness

https://arxiv.org/abs/2603.08309

202

Human-AI Divergence in Ego-centric Action Recognition under Spatial and Spatiotemporal Manipulations

https://arxiv.org/abs/2603.08317

203

Electrocardiogram Classification with Transformers Using Koopman and Wavelet Features

https://arxiv.org/abs/2603.08339

204

One Model Is Enough: Native Retrieval Embeddings from LLM Agent Hidden States

https://arxiv.org/abs/2603.08429

205

A prospective clinical feasibility study of a conversational diagnostic AI in an ambulatory primary care clinic

https://arxiv.org/abs/2603.08448

206

R2F: Repurposing Ray Frontiers for LLM-free Object Navigation

https://arxiv.org/abs/2603.08475

207

Visual Self-Fulfilling Alignment: Shaping Safety-Oriented Personas via Threat-Related Images

https://arxiv.org/abs/2603.08486

208

Towards Effective and Efficient Graph Alignment without Supervision

https://arxiv.org/abs/2603.08526

209

OSS-CRS: Liberating AIxCC Cyber Reasoning Systems for Real-World Open-Source Security

https://arxiv.org/abs/2603.08566

210

PostTrainBench: Can LLM Agents Automate LLM Post-Training?

https://arxiv.org/abs/2603.08640

211

A New Lower Bound for the Random Offerer Mechanism in Bilateral Trade using AI-Guided Evolutionary Search

https://arxiv.org/abs/2603.08679

212

Split Federated Learning Architectures for High-Accuracy and Low-Delay Model Training

https://arxiv.org/abs/2603.08687

213

An Embedding-based Approach to Inconsistency-tolerant Reasoning with Inconsistent Ontologies

https://arxiv.org/abs/2304.01664

214

From LLM Reasoning to Autonomous AI Agents: A Comprehensive Review

https://arxiv.org/abs/2504.19678

215

Precision Proactivity: Measuring Cognitive Load in Real-World AI-Assisted Work

https://arxiv.org/abs/2505.10742

216

MMTU: A Massive Multi-Task Table Understanding and Reasoning Benchmark

https://arxiv.org/abs/2506.05587

217

Let's Think in Two Steps: Mitigating Agreement Bias in MLLMs with Self-Grounded Verification

https://arxiv.org/abs/2507.11662

218

UniCast: A Unified Framework for Instance-Conditioned Multimodal Time-Series Forecasting

https://arxiv.org/abs/2508.11954

219

Synthetic Homes: An Accessible Multimodal Pipeline for Producing Residential Building Data with Generative AI

https://arxiv.org/abs/2509.09794

220

Mapping Overlaps in Benchmarks through Perplexity in the Wild

https://arxiv.org/abs/2509.23488

221

Your Agent May Misevolve: Emergent Risks in Self-evolving LLM Agents

https://arxiv.org/abs/2509.26354

222

Human-Centered LLM-Agent System for Detecting Anomalous Digital Asset Transactions

https://arxiv.org/abs/2510.20102

223

Jr. AI Scientist and Its Risk Report: Autonomous Scientific Exploration from a Baseline Paper

https://arxiv.org/abs/2511.04583

224

Parallel Decoder Transformer: Planner-Seeded Latent Coordination for Synchronized Parallel Decoding

https://arxiv.org/abs/2512.10054

225

Adaptation of Agentic AI: A Survey of Post-Training, Memory, and Skills

https://arxiv.org/abs/2512.16301

226

BoxMind: Closed-loop AI strategy optimization for elite boxing validated in the 2024 Olympics

https://arxiv.org/abs/2601.11492

227

Replayable Financial Agents: A Determinism-Faithfulness Assurance Harness for Tool-Using LLM Agents

https://arxiv.org/abs/2601.15322

228

MERIT Feedback Elicits Better Bargaining in LLM Negotiators

https://arxiv.org/abs/2602.10467

229

SkillsBench: Benchmarking How Well Agent Skills Work Across Diverse Tasks

https://arxiv.org/abs/2602.12670

230

Can a Lightweight Automated AI Pipeline Solve Research-Level Mathematical Problems?

https://arxiv.org/abs/2602.13695

231

Characterizing MARL for Energy Control: A Multi-KPI Benchmark on the CityLearn Environment

https://arxiv.org/abs/2602.19223

232

Vibe Researching as Wolf Coming: Can AI Agents with Skills Replace or Augment Social Scientists?

https://arxiv.org/abs/2602.22401

233

How Well Do Multimodal Models Reason on ECG Signals?

https://arxiv.org/abs/2603.00312

234

Extended Empirical Validation of the Explainability Solution Space

https://arxiv.org/abs/2603.01235

235

Opponent State Inference Under Partial Observability: An HMM-POMDP Framework for 2026 Formula 1 Energy Strategy

https://arxiv.org/abs/2603.01290

236

No Memorization, No Detection: Output Distribution-Based Contamination Detection in Small Language Models

https://arxiv.org/abs/2603.03203

237

Building Effective AI Coding Agents for the Terminal: Scaffolding, Harness, Context Engineering, and Lessons Learned

https://arxiv.org/abs/2603.05344

238

Exploring Diffusion Models' Corruption Stage in Few-Shot Fine-tuning and Mitigating with Bayesian Neural Networks

https://arxiv.org/abs/2405.19931

239

Multimodal Laryngoscopic Video Analysis for Assisted Diagnosis of Vocal Fold Paralysis

https://arxiv.org/abs/2409.03597

240

The Future of Software Testing: AI-Powered Test Case Generation and Validation

https://arxiv.org/abs/2409.05808

241

Reconsidering the energy efficiency of spiking neural networks

https://arxiv.org/abs/2409.08290

242

Rethinking SNN Online Training and Deployment: Gradient-Coherent Learning via Hybrid-Driven LIF Model

https://arxiv.org/abs/2410.07547

243

Efficient Semi-Supervised Adversarial Training via Latent Clustering-Based Data Reduction

https://arxiv.org/abs/2501.10466

244

Mitigating Unintended Memorization with LoRA in Federated Learning for LLMs

https://arxiv.org/abs/2502.05087

245

Prompt-SID: Learning Structural Representation Prompt via Latent Diffusion for Single-Image Denoising

https://arxiv.org/abs/2502.06432

246

Language in the Flow of Time: Time-Series-Paired Texts Weaved into a Unified Temporal Narrative

https://arxiv.org/abs/2502.08942

247

Enhancing Alzheimer's Diagnosis: Leveraging Anatomical Landmarks in Graph Convolutional Neural Networks on Tetrahedra...

https://arxiv.org/abs/2503.05031

248

From 2D Alignment to 3D Plausibility: Unifying Heterogeneous 2D Priors and Penetration-Free Diffusion for Occlusion-R...

https://arxiv.org/abs/2503.17788

249

MediTools -- Medical Education Powered by LLMs

https://arxiv.org/abs/2503.22769

250

"That's another doom I haven't thought about": A User Study on AI Labels as a Safeguard Against Image-Based Misinform...

https://arxiv.org/abs/2505.22845

251

Context Matters! Relaxing Goals with LLMs for Feasible 3D Scene Planning

https://arxiv.org/abs/2506.15828

252

Noisy PDE Training Requires Bigger PINNs

https://arxiv.org/abs/2507.06967

253

CauKer: Classification Time Series Foundation Models Can Be Pretrained on Synthetic Data

https://arxiv.org/abs/2508.02879

254

Entropy-Driven Curriculum for Multi-Task Training in Human Mobility Prediction

https://arxiv.org/abs/2509.01613

255

M4Diffuser: Multi-View Diffusion Policy with Manipulability-Aware Control for Robust Mobile Manipulation

https://arxiv.org/abs/2509.14980

256

Tiny but Mighty: A Software-Hardware Co-Design Approach for Efficient Multimodal Inference on Battery-Powered Small D...

https://arxiv.org/abs/2510.05109

257

Membership Inference Attacks on Tokenizers of Large Language Models

https://arxiv.org/abs/2510.05699

258

Bee: A High-Quality Corpus and Full-Stack Suite to Unlock Advanced Fully Open MLLMs

https://arxiv.org/abs/2510.13795

259

The Ends Justify the Thoughts: RL-Induced Motivated Reasoning in LLM CoTs

https://arxiv.org/abs/2510.17057

260

Explainable Heterogeneous Anomaly Detection in Financial Networks via Adaptive Expert Routing

https://arxiv.org/abs/2510.17088

261

Reinforcing Numerical Reasoning in LLMs for Tabular Prediction via Structural Priors

https://arxiv.org/abs/2510.17385

262

Vectorized Online POMDP Planning

https://arxiv.org/abs/2510.27191

263

FATE: A Formal Benchmark Series for Frontier Algebra of Multiple Difficulty Levels

https://arxiv.org/abs/2511.02872

264

Detecting AI-Generated Images via Contextual Anomaly Estimation in Masked AutoEncoders

https://arxiv.org/abs/2511.06325

265

Towards Realistic Guarantees: A Probabilistic Certificate for SmoothLLM

https://arxiv.org/abs/2511.18721

266

Automating Deception: Scalable Multi-Turn LLM Jailbreaks

https://arxiv.org/abs/2511.19517

267

ForamDeepSlice: A High-Accuracy Deep Learning Framework for Foraminifera Species Classification from 2D Micro-CT Slices

https://arxiv.org/abs/2512.00912

268

Beyond Additivity: Sparse Isotonic Shapley Regression toward Nonlinear Explainability

https://arxiv.org/abs/2512.03112

269

Beyond Endpoints: Path-Centric Reasoning for Vectorized Off-Road Network Extraction

https://arxiv.org/abs/2512.10416

270

The Algorithmic Gaze of Image Quality Assessment: An Audit and Trace Ethnography of the LAION-Aesthetics Predictor

https://arxiv.org/abs/2601.09896

271

Continuous-Flow Data-Rate-Aware CNN Inference on FPGA

https://arxiv.org/abs/2601.19940

272

MeanCache: From Instantaneous to Average Velocity for Accelerating Flow Matching Inference

https://arxiv.org/abs/2601.19961

273

Impact of LLMs news Sentiment Analysis on Stock Price Movement Prediction

https://arxiv.org/abs/2602.00086

274

Thickening-to-Thinning: Reward Shaping via Human-Inspired Learning Dynamics for LLM Reasoning

https://arxiv.org/abs/2602.04265

275

Extracting Recurring Vulnerabilities from Black-Box LLM-Generated Software

https://arxiv.org/abs/2602.04894

276

LMMRec: LLM-driven Motivation-aware Multimodal Recommendation

https://arxiv.org/abs/2602.05474

277

Explainable Token-level Noise Filtering for LLM Fine-tuning Datasets

https://arxiv.org/abs/2602.14536

278

LongAudio-RAG: Event-Grounded Question Answering over Multi-Hour Long Audio

https://arxiv.org/abs/2602.14612

279

Understand Then Memory: A Cognitive Gist-Driven RAG Framework with Global Semantic Diffusion

https://arxiv.org/abs/2602.15895

280

CrystaL: Spontaneous Emergence of Visual Latents in MLLMs

https://arxiv.org/abs/2602.20980

281

MrBERT: Modern Multilingual Encoders via Vocabulary, Domain, and Dimensional Adaptation

https://arxiv.org/abs/2602.21379

282

CryoNet.Refine: A One-step Diffusion Model for Rapid Refinement of Structural Models with Cryo-EM Density Map Restraints

https://arxiv.org/abs/2602.22263

283

Attn-QAT: 4-Bit Attention With Quantization-Aware Training

https://arxiv.org/abs/2603.00040

284

Conformal Prediction for Risk-Controlled Medical Entity Extraction Across Clinical Domains

https://arxiv.org/abs/2603.00924

285

Slurry-as-a-Service: A Modest Proposal on Scalable Pluralistic Alignment for Nutrient Optimization

https://arxiv.org/abs/2603.02420

286

Human-Certified Module Repositories for the AI Age

https://arxiv.org/abs/2603.02512

287

ITO: Images and Texts as One via Synergizing Multiple Alignment and Training-Time Fusion

https://arxiv.org/abs/2603.02767

288

Interpretable Motion-Attentive Maps: Spatio-Temporally Localizing Concepts in Video Diffusion Transformers

https://arxiv.org/abs/2603.02919

289

Why Adam Can Beat SGD: Second-Moment Normalization Yields Sharper Tails

https://arxiv.org/abs/2603.03099

290

Information Routing in Atomistic Foundation Models: How Task Alignment and Equivariance Shape Linear Disentanglement

https://arxiv.org/abs/2603.03155

291

SAIL: Similarity-Aware Guidance and Inter-Caption Augmentation-based Learning for Weakly-Supervised Dense Video Capti...

https://arxiv.org/abs/2603.05437

292

When AI Levels the Playing Field: Skill Homogenization, Asset Concentration, and Two Regimes of Inequality

https://arxiv.org/abs/2603.05565

293

Bridging Domains through Subspace-Aware Model Merging

https://arxiv.org/abs/2603.05768